

  Kohler

Tuesday 12:33 AM

KOHLER AI Research Lab — Phase 1: Enterprise Research.
Continue from Track 3 — Unified Enterprise AI Agent.

Absolutely. We'll continue with **Track 3 — Unified Enterprise AI Agent** and start **Phase 1: Enterprise Research**.

The important framing we already established is:

We are not building a generic chatbot.

We are researching and designing an enterprise AI agent capable of orchestrating KOHLER's fragmented **knowledge, enterprise data, tools, workflows, and specialist agents**, while respecting permissions, governance, human approval, and observability.

Phase 1 — Enterprise Research

The goal of Phase 1 is to understand what KOHLER actually looks like as an enterprise before deciding what our agent should build.

We'll research it in **5 areas**:

1. KOHLER Enterprise Landscape

Understand:

- KOHLER's business divisions
- Major products and services
- Global operations
- Manufacturing and supply chain
- Sales and distribution
- Corporate functions
- Digital/technology organization
- Enterprise software environment

Question we need to answer:

Where does information and work currently live inside KOHLER?

2. KOHLER's Existing AI Ecosystem

We already have an important starting point from the earlier research:

- 98% Copilot adoption
- Around 3,000 agents

Rather than treating these as just impressive numbers, we'll investigate:

If KOHLER already has thousands of AI agents, why would a Unified Enterprise AI Agent be valuable?

This is probably one of the most important questions in our entire case study.

We'll investigate:


- Microsoft Copilot usage
- Existing AI agents
- AI initiatives
- Agentic AI strategy
- Internal AI platforms
- AI use cases
- Employee adoption
- Existing AI governance

3. Enterprise Pain Points

This is where we'll identify the actual problem our system should solve.

We'll investigate problems such as:

Information fragmentation



↓
Existing AI agents

An employee may need information from several of these systems to complete **one task**.

So the problem isn't necessarily:

"KOHLER doesn't have AI."

It may instead be:

"KOHLER has many AI capabilities, but employees still need to know where to go, which agent to use, which system contains the required information, and how to combine the results."

That distinction will be central to Track 3.

4. Enterprise Agent Architecture Research

Once we understand KOHLER, we'll research how a unified enterprise agent could technically work.

We'll investigate:

Agent orchestration

How does a central agent decide:

```
graph TD
    A[User request] --> B[Understand intent]
    B --> C[Plan task]
    C --> D[Select tools / agents]
    D --> E[Retrieve enterprise data]
    E --> F[Reason]
    F --> G[Execute actions]
    G --> H[Verify result]
    H --> I[Respond]
```

Enterprise RAG

We'll research:

- document retrieval
- vector databases
- hybrid search
- metadata filtering
- permission-aware retrieval
- enterprise knowledge graphs


Tool use

The agent may need to interact with:

ERP
CRM
HR systems
Manufacturing systems
Databases
Email
Calendar
Document repositories
Existing AI agents
Internal APIs

Multi-agent orchestration

Instead of one giant agent:



We'll investigate whether this architecture actually makes sense for KOHLER.

5. Enterprise Security & Governance

This is **critical**.

A consumer chatbot can potentially answer:

"What is our revenue?"


An enterprise agent needs to ask:

"Is this particular employee authorized to see revenue information?"

So we'll research:

- Identity & access management
- RBAC / ABAC
- Permission-aware RAG
- Data classification
- Data privacy
- Audit logs
- Human-in-the-loop
- Approval workflows
- Prompt injection
- Data leakage
- Agent permissions
- Tool authorization
- Enterprise AI governance
- Observability

The architecture should therefore look more like:





Our Phase 1 Research Questions

Rather than randomly researching KOHLER, we'll answer a defined set of questions.

#	Research Question
1	What does KOHLER's enterprise structure look like?
2	What are KOHLER's major business functions?
3	What technology/data systems are likely involved?
4	What AI initiatives has KOHLER already implemented?
5	What does the 98% Copilot adoption actually mean?
6	What are the ~3,000 agents and what role do they play?
7	What enterprise problems remain despite widespread AI adoption?
8	Where is information/work fragmented?
9	Which employee workflows could benefit from orchestration?
10	What should a unified enterprise agent actually control?
11	What security/governance constraints must it respect?
12	How should we evaluate whether the proposed agent is successful?

Phase 1 Deliverable

At the end of this phase, we should have a proper **Enterprise Research Brief**, not just notes.

It will contain:

A. KOHLER Profile

Business + organizational landscape

B. Existing AI Landscape

Copilot + agents + AI initiatives

C. Enterprise Technology Landscape

Systems, data and knowledge sources

D. Pain-Point Analysis

Where fragmentation creates friction

E. Opportunity Analysis

Where a unified agent creates value

F. Target Users

Which employees/functions benefit

G. Initial Use Cases

Concrete workflows for our prototype

H. Technical Requirements

What the architecture must support

I. Security & Governance Requirements

J. Research Conclusions

The exact problem statement we'll take into Phase 2.

One important rule for our research

We'll separate **facts** from **assumptions**.

For example:

Fact

KOHLER publicly states X.

Inference

This suggests KOHLER may have Y requirement.

Our design decision

Therefore, our prototype will implement Z.

That will make the final case study much stronger because we're not pretending to know KOHLER's internal architecture when we don't have public evidence.

Let's start with Research Area 1

Next: KOHLER Enterprise Landscape.

We'll research KOHLER's:

Business → Organization → Operations → Technology → AI

and build our first evidence-backed picture of the enterprise.

Then we'll move to **KOHLER's existing AI ecosystem**, where the **98% Copilot adoption / ~3,000 agents** becomes especially important.

      Sources

Area 4 — What Should We Build?

Important Takeaways

Our solution is a **Unified Enterprise AI Agent** that acts as an **orchestration layer** over KOHLER's existing AI ecosystem.

It should:

1. **Understand user intent** — identify what the employee actually needs.
2. **Search enterprise knowledge** — retrieve relevant authorized information.
3. **Discover agents** — find the appropriate specialized agents.
4. **Orchestrate agents** — coordinate multiple agents for complex tasks.
5. **Use tools/APIs** — interact with enterprise systems.
6. **Reason & synthesize** — combine results into one useful answer.
7. **Execute with approval** — perform actions while keeping humans in control.



Notes to Write

Area 4 — Solution

- Build a **Unified Enterprise AI Agent**.
- Acts as an **orchestration layer**, not a replacement for existing agents.
- Single interface for employees.
- Connects:
 - Existing AI agents
 - Enterprise data
 - Business tools/APIs
- Core capabilities:
Intent → Retrieve → Select → Orchestrate → Execute → Verify → Respond
- High-impact actions require **human approval**.


Main Prototype Use Case

Cross-functional Order Investigation

Example:

"Why is Order #12345 delayed? Find the reason, check the impact, and prepare a

customer response."



Key Insight

Don't replace KOHLER's existing agents. Build a unified layer that makes those agents, data and tools work together.



area 5

Area 5 — Security & Governance

Important Takeaways

Because our agent can access enterprise data and perform actions, **security and governance are essential.**

1. Identity & Access Control

- Identify the user.
- Only provide data the user is authorized to access.

2. Permission-Aware RAG

- Filter documents/data based on user permissions **before retrieval.**
- The AI must not bypass existing access controls.

3. Least Privilege

- Each agent/tool gets only the permissions it needs.

4. Human-in-the-Loop

- Low-risk → AI can answer directly.
- High-impact → require human approval before execution.

5. Auditability

- Log important actions:

User → Request → Data → Agent → Tool → Action → Approval

6. AI Security

- Protect against prompt injection, data leakage and unauthorized tool usage.

Notes to Write

Area 5 — Security & Governance

- **Identity & Access:** Verify who the user is.
- **Permission-Aware RAG:** Retrieve only authorized data.
- **Least Privilege:** Give agents minimum required permissions.
- **Human Approval:** Required for sensitive/high-impact actions.
- **Audit Logs:** Track users, agents, data, tools and actions.
- **AI Security:** Protect against prompt injection and data leakage.
- **Fail Safely:** If uncertain or unauthorized → don't act; ask/escalate.

Key Insight

The Unified Agent must be powerful enough to orchestrate enterprise systems but restricted enough that it cannot bypass permissions or take uncontrolled actions.

Phase 1 — One-Page Summary

You now have the complete research:

Area	Core Idea
1. Enterprise Landscape	KOHLER is a large, global, diversified enterprise.
2. Existing AI	98% Copilot adoption + ~3,000 agents.
3. Pain Points	Discovery, fragmented data, coordination, execution, context switching.
4. Solution	Unified orchestration layer connecting agents, data & tools.


Area

5. Security

Core Idea

Permissions, least privilege, human approval & auditability.

The whole Track 3 in one diagram:



Phase 1 is now complete. We can move to Phase 2 — Solution & Architecture Design, where we'll turn this research into the actual system we're going to propose/build.



lets do the phase 2 in next chat

Absolutely. We'll continue in a new chat with Phase 2 — Solution & Architecture Design.

When you start it, just say:

"Continue KOHLER AI Research Lab — Track 3, Phase 2."

We'll pick up directly from the Phase 1 conclusions and start designing the **Unified Enterprise AI Agent**—architecture, agents, RAG, tools/APIs, orchestration, memory,

security, and the prototype scope.



ChatGPT can make mistakes. Check important info.